

## diabetes prediction using logistic regression

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: ds = pd.read_csv('diabetes.csv')
```


```
In [3]: ds.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                        768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [6]: ds.describe()
```

```
Out[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	33.240923	33.868129	0.345581
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	33.973990	5.415014	0.476946
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	33.000000	24.000000	0.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	33.000000	33.000000	0.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	33.000000	41.000000	0.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	641.940000	50.000000	1.000000



```
In [7]: ds.isnull().sum()
```


```
Out[7]: Pregnancies      0
        Glucose          0
        BloodPressure    0
        SkinThickness    0
        Insulin          0
        BMI              0
        DiabetesPedigreeFunction 0
        Age              0
        Outcome          0
        dtype: int64
```

```
In [8]: ds.duplicated().sum()
```

```
Out[8]: np.int64(0)
```

```
In [9]: sns.countplot(x="Outcome", data=ds)
```


```
Out[9]: <Axes: xlabel='Outcome', ylabel='count'>
```



```
In [10]: sns.pairplot(data=ds, hue="Outcome")
         plt.show()
```



```
In [11]: sns.heatmap(ds.corr(),annot = True)
plt.show()
```



```
In [13]: ds_new = ds
ds_new[["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]] = ds_new[["
```

```
ds_new["Glucose"].fillna(ds_new["Glucose"].mean() , inplace = True)
ds_new["BloodPressure"].fillna(ds_new["BloodPressure"].mean() , inplace = True)
ds_new["Insulin"].fillna(ds_new["Insulin"].mean() , inplace = True)
ds_new["BMI"].fillna(ds_new["BMI"].mean() , inplace = True)
ds_new["SkinThickness"].fillna(ds_new["SkinThickness"].mean() , inplace = True)
```

C:\Users\Aryansh Pathak\AppData\Local\Temp\ipykernel\_13604\1196353376.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
ds_new["Glucose"].fillna(ds_new["Glucose"].mean() , inplace = True)
```

C:\Users\Aryansh Pathak\AppData\Local\Temp\ipykernel\_13604\1196353376.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
ds_new["BloodPressure"].fillna(ds_new["BloodPressure"].mean() , inplace = True)
```

C:\Users\Aryansh Pathak\AppData\Local\Temp\ipykernel\_13604\1196353376.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
ds_new["Insulin"].fillna(ds_new["Insulin"].mean() , inplace = True)
```

C:\Users\Aryansh Pathak\AppData\Local\Temp\ipykernel\_13604\1196353376.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
ds_new["BMI"].fillna(ds_new["BMI"].mean() , inplace = True)
```

C:\Users\Aryansh Pathak\AppData\Local\Temp\ipykernel\_13604\1196353376.py:5: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform



the operation inplace on the original object.

```
ds_new["SkinThickness"].fillna(ds_new["SkinThickness"].mean() , inplace = True)
```

```
In [15]: ds_new.isnull().sum()
```

```
Out[15]: Pregnancies      0
          Glucose          0
          BloodPressure    0
          SkinThickness     0
          Insulin           0
          BMI               0
          DiabetesPedigreeFunction  0
          Age              0
          Outcome          0
          dtype: int64
```

```
In [16]: y = ds_new['Outcome']
          x = ds_new.drop('Outcome' , axis = 1)
```

```
In [17]: from sklearn.model_selection import train_test_split
          x_train , x_test , y_train , y_test = train_test_split(x, y , test_size = 0.20 , ra
```

```
In [18]: from sklearn.linear_model import LogisticRegression
          model = LogisticRegression()
          model.fit(x_train , y_train)
          y_predict = model.predict(x_test)
```

C:\Users\Aryansh Pathak\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\linear\_model\\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

```
In [19]: y_predict
```


```
Out[19]: array([1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
                0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
                0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0,
                1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0,
                1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1,
                0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0])
```

```
In [20]: from sklearn.metrics import confusion_matrix
          cm = confusion_matrix(y_test , y_predict)
          cm
```

```
Out[20]: array([[82, 18],
                [28, 26]])
```

```
In [21]: sns.heatmap(pd.DataFrame(cm), annot=True)
```

```
Out[21]: <Axes: >
```



```
In [22]: from sklearn.metrics import accuracy_score
```

```
In [23]: accuracy = accuracy_score(y_test,y_predict)
accuracy
```

```
Out[23]: 0.7012987012987013
```

```
In [24]: y_predict = model.predict([[1,148,72,38,71.789,13.6,0.927,10]])
print(y_predict)
if y_predict==1:
    print("Diabetic")
else:
    print("Non Diabetic")
```

```
[0]
Non Diabetic
```

```
C:\Users\Aryansh Pathak\AppData\Local\Programs\Python\Python312\Lib\site-packages\sk
learn\base.py:493: UserWarning: X does not have valid feature names, but LogisticReg
ression was fitted with feature names
  warnings.warn(
```

```
In [ ]:
```